

Blind Refusal: Language Models Refuse to Help Users Evade Unjust, Absurd, and Illegitimate Rules

Cameron Pattison¹ Lorenzo Manuali² Seth Lazar³


¹Department of Philosophy, Vanderbilt University ²Department of Philosophy, University of Michigan ³School of Government and Policy, Johns Hopkins University

The failure mode

Safety-trained models refuse to help users circumvent rules. But not all rules deserve compliance.

Blind refusal: the model refuses without asking whether the rule is just, the authority legitimate, or the application fair.


Overrefusal benchmarks ask whether the request is harmful. We ask whether the rule deserves enforcement.



Verbatim responses from dataset. Scenario: journalist under authoritarian censorship (seed_illegitimate_authority_national_government_22). Both models received the identical prompt.

Fig. 1. Same request, two chatbots: one helps, one refuses.

Design



Results

Models refuse to help with defeated rules in three out of four cases

75.4% refusal on defeated rules vs 95.8% on control, and 59.6% even on cases the dual-use gate marks safe.




Fig. 2. Refusal rate by defeat type × authority type, averaged over 18 configurations. Control (top row) is near-universal; defeated rows stay high.




Fig. 3. Refusal by defeat family, per configuration. Larger profile = more refusal. GPT-5.4 variants are most inflated; Grok-4 is small everywhere, including control.

Key numbers

Case-clustered bootstrap, 95% CI	
Refusal, defeated rules	75.4% [73.8, 77.1]
Refusal, control	95.8% [94.8, 96.7]
Helps, defeated – control	+20.3 pp [18.5, 22.2]
on non-dual-use cases	+33.4 pp
on dual-use cases	+9.2 pp
Refusal on gate-safe cases	59.6%
Engages with the defeat	57.5% vs 4.5% control
Refusal when engaged	74.1% vs 77.2% unengaged

N = 23,189 evaluations (1,290 cases × 18 configurations; 31 pairs dropped). Mixed-effects model: engaging with the defeat raises the odds of helping, OR 1.77 [1.56, 2.02].

Two ways to fail

Help on defeated rules ranges from 7.7% (GPT-5.4-mini-base) to 58.0% (Grok-4).

GPT-5.4 refuses >89% of defeated-rule requests; reasoning changes nothing (9.7% vs 9.9% helps).

Grok-4 helps with 58% of defeated cases, but also 33.5% of control: permissive, not discriminating.

Takeaways

- Models see that the rule is questionable and refuse anyway (74.1% refusal when engaged vs 77.2% when not). The failure sits downstream of content classification.
- More helpfulness is not the fix. Grok-4 shows what indiscriminate permissiveness looks like.
- Alignment needs sensitivity to when rules may legitimately be broken, a question political philosophy has long theorized and AI safety has yet to operationalize.

Limitations. The dual-use gate flags 57% of cases and may over-flag, which would inflate the non-dual-use delta. The independent-harm flag fires on roughly 70% of non-help responses, likely too often, which if anything understates blind refusal. Only obvious injustices pass the reasonable-judge gate; contested cases remain unmeasured.

Acknowledgments. Thank you to Angelica Chowdhury for substantial assistance in our statistical analysis, and to the MINT Lab for audience and early feedback on previous versions of this project.



Code and data
github.com/mint-philosophy/blind-refusal



COLM 2026